

# Distributed deductive databases, declaratively: The L10 logic programming language

Robert J. Simmons  
Bernardo Toninho  
Frank Pfenning

Computer Science Department  
Carnegie Mellon University

June 4, 2011

# Summary

- 1 Introduction
- 2 L10 Language Features
- 3 On Implementing L10
- 4 Future Work and Conclusion

# Introduction

## What is L10?

### What is L10?

- *Forward-chaining* logic programming language,
- Distribution of data and parallelism of computation,
- Logically motivated notion of worlds as locations for computation.

### Forward-chaining logic programming

- Deals with collections of facts that model some structure;
- Computation is described by rules;
- Operational interpretation of rules is exhaustive forward deduction:
  - Try to match facts from the database against the premises of a rule;
  - Add the conclusion to the database, if not already present;
  - Computation terminates when no new facts can be added.

# Introduction

## What is L10?

### What is L10?

- *Forward-chaining* logic programming language,
- Distribution of data and parallelism of computation,
- Logically motivated notion of worlds as locations for computation.

### Forward-chaining logic programming


- Deals with collections of facts that model some structure;
- Computation is described by rules;
- Operational interpretation of rules is exhaustive forward deduction:
  - Try to match facts from the database against the premises of a rule;
  - Add the conclusion to the database, if not already present;
  - Computation terminates when no new facts can be added.

# Introduction

## What is L10? - Forward-chaining

### Transitive closure of a graph

We can encode a graph as a fact database:


$$\left\{ \begin{array}{lll} \text{edge } a \ c & \text{edge } c \ a & \text{edge } d \ c \\ \text{edge } a \ d & \text{edge } c \ d & \text{edge } d \ d \\ \text{edge } b \ b & \text{edge } d \ b & \end{array} \right\}$$

`edge X Y -> path X Y.`

`edge X Y, path Y Z -> path X Z.`

### Why forward-chaining?


- A natural way to describe fixed-point iteration and database-like algorithms;
- Can produce surprisingly succinct and efficient programs.

# Introduction

## What is L10? - Forward-chaining

### Transitive closure of a graph

We can encode a graph as a fact database:


$$\left\{ \begin{array}{lll} \text{edge } a \ c & \text{edge } c \ a & \text{edge } d \ c \\ \text{edge } a \ d & \text{edge } c \ d & \text{edge } d \ d \\ \text{edge } b \ b & \text{edge } d \ b & \end{array} \right\}$$

edge  $X \ Y \rightarrow$  path  $X \ Y$ .

edge  $X \ Y$ , path  $Y \ Z \rightarrow$  path  $X \ Z$ .

### Why forward-chaining?


- A natural way to describe fixed-point iteration and database-like algorithms;
- Can produce surprisingly succinct and efficient programs.

# Introduction

## What is L10? - Forward-chaining

### Transitive closure of a graph

We can encode a graph as a fact database:


$$\left\{ \begin{array}{lll} \text{edge } a \ c & \text{edge } c \ a & \text{edge } d \ c \\ \text{edge } a \ d & \text{edge } c \ d & \text{edge } d \ d \\ \text{edge } b \ b & \text{edge } d \ b & \end{array} \right\}$$

edge  $X \ Y \rightarrow$  path  $X \ Y$ .

edge  $X \ Y$ , path  $Y \ Z \rightarrow$  path  $X \ Z$ .

### Why forward-chaining?

- A natural way to describe fixed-point iteration and database-like algorithms;
- Can produce surprisingly succinct and efficient programs.

# Summary

- 1 Introduction
- 2 L10 Language Features**
- 3 On Implementing L10
- 4 Future Work and Conclusion



# L10 Language Features

## Distribution

### Distributed programming in L10

The foundation of distributed programming in L10 is *worlds*:

- Abstractly represent different storage and computation locations;
- All relations in L10 must be associated with a declared world;
- Dependencies between relations result in dependencies between worlds.

### Example: Liveness analysis declaration

L10 is a typed language. All worlds and relations must be declared:

```
wLive : world.  
live : nat -> t -> rel @ wLive.
```

# L10 Language Features

## Distribution

### Distributed programming in L10

The foundation of distributed programming in L10 is *worlds*:

- Abstractly represent different storage and computation locations;
- All relations in L10 must be associated with a declared world;
- Dependencies between relations result in dependencies between worlds.

### Example: Liveness analysis declaration

L10 is a typed language. All worlds and relations must be declared:

```
wLive : world.  
live  : nat -> t -> rel @ wLive.
```

# L10 Language Features

## Constructive Negation

L10 worlds stage computation by determining the order in which relations are computed.

### Program Analyses - Liveness

Information about the program code is encoded using these relations:

- `def L X` - `X` is defined in line `L`.
- `use L X` - `X` is used in line `L`.
- `succ L L'` - `L'` may be executed immediately after `L`.

Liveness is defined by the two rules:

`use L X -> live L X.`

`live L' X, succ L L', not (def L X) -> live L X.`

Negation in forward-chaining logic programming can be problematic...

# L10 Language Features

## Constructive Negation

### Stratified Negation

Some uses of negation can make sense. For instance:

`not (fact2) -> fact1`

where we can stage computation such that `fact2` is *completely* determined when we are considering this rule for `fact1` .

```
wCode : world.  
def   : nat -> t -> rel @ wCode.  
...  
live : nat -> t -> rel @ wLive.  
use L X -> live L X.  
live L' X, succ L L', not (def L X) -> live L X.
```

No *cyclic* dependencies between worlds are allowed in L10.

# L10 Language Features

## Constructive Negation

### Stratified Negation

Some uses of negation can make sense. For instance:

```
not (fact2) -> fact1
```

where we can stage computation such that `fact2` is *completely* determined when we are considering this rule for `fact1`.

```
wCode : world.  
def   : nat -> t -> rel @ wCode.  
...  
live : nat -> t -> rel @ wLive.  
use L X -> live L X.  
live L' X, succ L L', not (def L X) -> live L X.
```

No *cyclic* dependencies between worlds are allowed in L10.

# L10 Language Features

## Parallelism

### Exploiting worlds for parallelism

We can safely stage independent worlds for parallel execution.

### Program Analyses - Neededness

We can define a neededness analysis:

- `nec L X @ wCode`: at line L, X is necessary for control flow or as the return value.
- `needed L X @ wNeed`: at line L, X is needed.

`nec L X -> needed L X.`

`needed L' X, succ L L', not (def L X) -> needed L X.`

`use L Y, def L X, succ L L', needed L' X -> needed L Y.`

This way, the liveness and neededness analyses can be executed in parallel.

# L10 Language Features

## Parallelism

### Exploiting worlds for parallelism

We can safely stage independent worlds for parallel execution.

### Program Analyses - Neededness

We can define a neededness analysis:

- `nec L X @ wCode`: at line L, X is necessary for control flow or as the return value.
- `needed L X @ wNeed`: at line L, X is needed.

`nec L X -> needed L X.`

`needed L' X, succ L L', not (def L X) -> needed L X.`

`use L Y, def L X, succ L L', needed L' X -> needed L Y.`

This way, the liveness and neededness analyses can be executed in parallel.

# L10 Language Features

Indexed worlds and limited saturation

Suppose we want to implement a regular expression matcher. The type `regexp` captures the form of reg. expressions. Tokens will be represented by string constants.

## Reg. exp. matcher declaration

```
regexp : type.  
tok : string -> regexp.  
emp : regexp.  
alt : regexp -> regexp -> regexp.  
neg : regexp -> regexp.
```

```
w0 : world.  
w1 : regexp -> world.  
token : string -> nat -> rel @ w0.  
match : {RE : regexp} nat -> nat -> rel @ w1 RE.
```



# L10 Language Features

Indexed worlds and limited saturation

Suppose we want to implement a regular expression matcher. The type `regexp` captures the form of reg. expressions. Tokens will be represented by string constants.

## Reg. exp. matcher declaration

```
regexp : type.  
tok : string -> regexp.  
emp : regexp.  
alt : regexp -> regexp -> regexp.  
neg : regexp -> regexp.
```

```
w0 : world.  
w1 : regexp -> world.  
token : string -> nat -> rel @ w0.  
match : {RE : regexp} nat -> nat -> rel @ w1 RE.
```

# L10 Language Features

Indexed worlds and limited saturation

```
...  
token _ I -> match emp I I.  
  
match RE1 I J -> match (alt RE1 RE2) I J.  
match RE2 I J -> match (alt RE1 RE2) I J.
```

Most deductive databases would not allow these rules for alternation!

```
token _ I, token _ J, I <= J,  
  not (match RE I J) -> match (neg RE) I J.
```

Negation is justified by *locally stratified negation*.

# L10 Language Features

Indexed worlds and limited saturation

```
...  
token _ I -> match emp I I.  
  
match RE1 I J -> match (alt RE1 RE2) I J.  
match RE2 I J -> match (alt RE1 RE2) I J.
```

Most deductive databases would not allow these rules for alternation!

```
token _ I, token _ J, I <= J,  
  not (match RE I J) -> match (neg RE) I J.
```

Negation is justified by *locally stratified negation*.

# L10 Language Features

Indexed worlds and limited saturation

```
...  
token _ I -> match emp I I.  
  
match RE1 I J -> match (alt RE1 RE2) I J.  
match RE2 I J -> match (alt RE1 RE2) I J.
```

Most deductive databases would not allow these rules for alternation!

```
token _ I, token _ J, I <= J,  
    not (match RE I J) -> match (neg RE) I J.
```

Negation is justified by *locally stratified negation*.

# Summary

- 1 Introduction
- 2 L10 Language Features
- 3 On Implementing L10**
- 4 Future Work and Conclusion

# On Implementing L10

## Scheduling

### Static scheduling

How to handle queries for non-indexed worlds:

- 1 Compute the world dependency graph.
- 2 Perform a breadth-first traversal of the graph.
- 3 Produce a task list that maps L10 worlds to available X10 places.

### Scheduling program analyses

Assuming two X10 places, A and B:

- 1 Worlds  $w_{Live}$  and  $w_{Need}$  depend on world  $w_{Code}$ , which depends on no worlds.
- 2 BFS traversal:  $w_{Code}$ , followed by  $w_{Live}$  and  $w_{Need}$ .
- 3 Schedule  $w_{Code}$  and  $w_{Live}$  at place A,  $w_{Need}$  at place B.  
Computation at place B blocks until relations at  $w_{Code}$  are computed.

# On Implementing L10

## Scheduling

### Static scheduling

How to handle queries for non-indexed worlds:

- 1 Compute the world dependency graph.
- 2 Perform a breadth-first traversal of the graph.
- 3 Produce a task list that maps L10 worlds to available X10 places.

### Scheduling program analyses

Assuming two X10 places, A and B:

- 1 Worlds  $w_{Live}$  and  $w_{Need}$  depend on world  $w_{Code}$ , which depends on no worlds.
- 2 BFS traversal:  $w_{Code}$ , followed by  $w_{Live}$  and  $w_{Need}$ .
- 3 Schedule  $w_{Code}$  and  $w_{Live}$  at place A,  $w_{Need}$  at place B.  
Computation at place B blocks until relations at  $w_{Code}$  are computed.

# On Implementing L10

## Scheduling indexed worlds


### Scheduling indexed worlds

Indexed worlds are more interesting:

- Perform a BF traversal of the relevant subterm indices of the world:
  - until all subterms have been considered,
  - or the number of unique branches exceeds available parallelism.

### Scheduling a regular expression match

Assuming 3 X10 places A, B and C, a matching for  $\neg(b \mid c)(a+)$  is scheduled as:





# On Implementing L10

## Status of Implementation

### Integration with X10

We do not specify how to query saturated databases:

- Such queries will be performed through an API within X10
- Main uses of L10 programs are to provide data to other programs
- The language will eventually be available through an X10 library
- Similar APIs exist for many deductive database/programming language combinations

L10 is still at a very early development stage:

- Fully functional interpreter (written in Standard ML) - fully sequential.
- Compiler to Standard ML and X10 - in development.

# Summary

- 1 Introduction
- 2 L10 Language Features
- 3 On Implementing L10
- 4 Future Work and Conclusion**

# Future Work & Conclusions

## Future Work

- Indexing worlds with non-structured terms (e.g. strings, numbers)
- Optimizing communication between worlds/places
- Finishing the implementation of the compiler...

## Conclusions

- Introduced the preliminary design of a rich distributed logic programming language
- Exploit a mapping of the logically motivated notion of worlds to X10 places,
- Take advantage of inherent parallelism present in logic programs

# Future Work & Conclusions

## Future Work

- Indexing worlds with non-structured terms (e.g. strings, numbers)
- Optimizing communication between worlds/places
- Finishing the implementation of the compiler...

## Conclusions

- Introduced the preliminary design of a rich distributed logic programming language
- Exploit a mapping of the logically motivated notion of worlds to X10 places,
- Take advantage of inherent parallelism present in logic programs

Suggestions are welcomed...