



Phaser Beams: Integrating Stream Parallelism with Task Parallelism

X10 Workshop

June 4th, 2011

Jun Shirako, David M. Peixotto,
Dragos-Dumitru Sbirlea and Vivek Sarkar
Rice University

Introduction

- **Stream Languages**

- Natures to explicitly specify streaming parallelism in a stream graph
 - Filter (node): Computation unit
 - Stream (edge): Flow of data among filters
- Lack of dynamic parallelism
 - Fixed stream graphs w/o dynamic reconfiguration

- **Task Parallel Languages**

- Support of dynamic task parallelism
 - Task: Dynamically created/terminated lightweight thread
 - Chapel, Cilk, Fortress, Habanero-Java/C, Intel Threading Building Blocks, Java Concurrency Utilities, Microsoft Task Parallel Library, OpenMP 3.0 and X10
- Lack of support for efficient streaming communication among tasks

- **Address the gap between two paradigms**

- Phaser beams: Integrating Stream and Dynamic Task parallel models

Introduction

- **Habanero-Java**
 - Task parallel language based on X10 v1.5
 - <http://habanero.rice.edu/hj>
- **Phasers in HJ**
 - Extension of X10 clocks
 - Synchronization for dynamic task parallel model
 - Various synchronization patterns
 - Collective barriers, point-to-point synchronizations
 - Java 7 Phasers
- **Streaming extensions to phasers**
 - Streaming communication among tasks
 - Adaptive batch optimization
 - Runtime cycle detection for efficient execution of acyclic stream graphs

Outline

- Introduction
- **Habanero-Java parallel constructs**
 - async, finish, phasers and accumulator
- **Extensions for streaming with dynamic parallelism**
 - Phaser beams
 - Expressed streaming patterns
- **Adaptive batch optimization**
 - Runtime cycle detection
 - Adaptive batching to avoid deadlock
- **Experimental results**
- **Conclusions**

Task Creation & Termination

- Async: Lightweight task creation
- Finish: Task-set termination

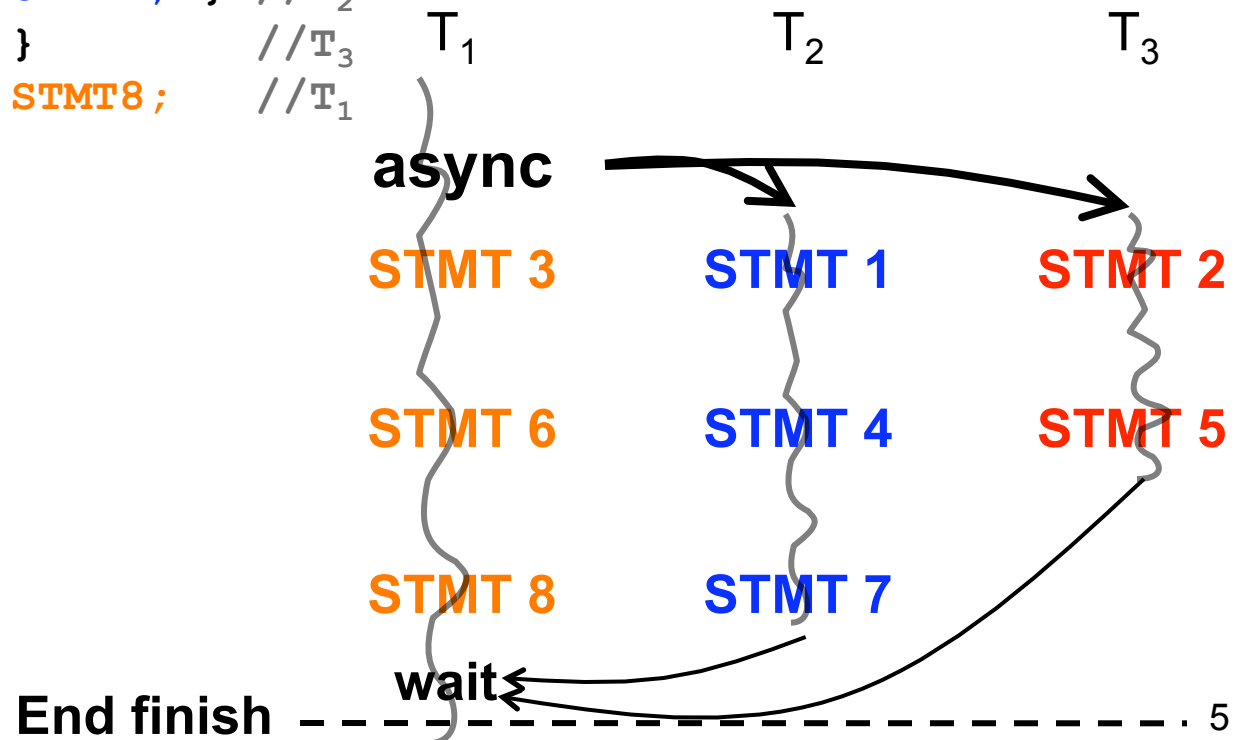
```
finish { // Start finish
```

```
    // T1 creates T2 and T3
```

```
    async { STMT1; STMT4; STMT7; } //T2
```

```
    async { STMT2; STMT5; } //T3  
           STMT3; STMT6; STMT8; //T1
```

```
} // End finish
```



Phasers

- **Phaser allocation**

- **phaser** **ph** = **new phaser(mode)**

- Phaser **ph** is allocated with **registration mode**
 - Mode: **SIG_WAIT_SINGLE** (default)



- Registration mode defines capability
 - There is a lattice ordering of capabilities

- **Task registration**

- **async phased** (**ph₁**<**mode₁**>, **ph₂**<**mode₂**>, ...) {**STMT**}

- Created task is registered with **ph₁** in **mode₁**, **ph₂** in **mode₂**, ...
 - **Capability rule**: Child task's registration mode must be subset of parent's

- **Synchronization**

- **next**: Equivalent to **signal** followed by **wait**

- Deadlock-free execution semantics

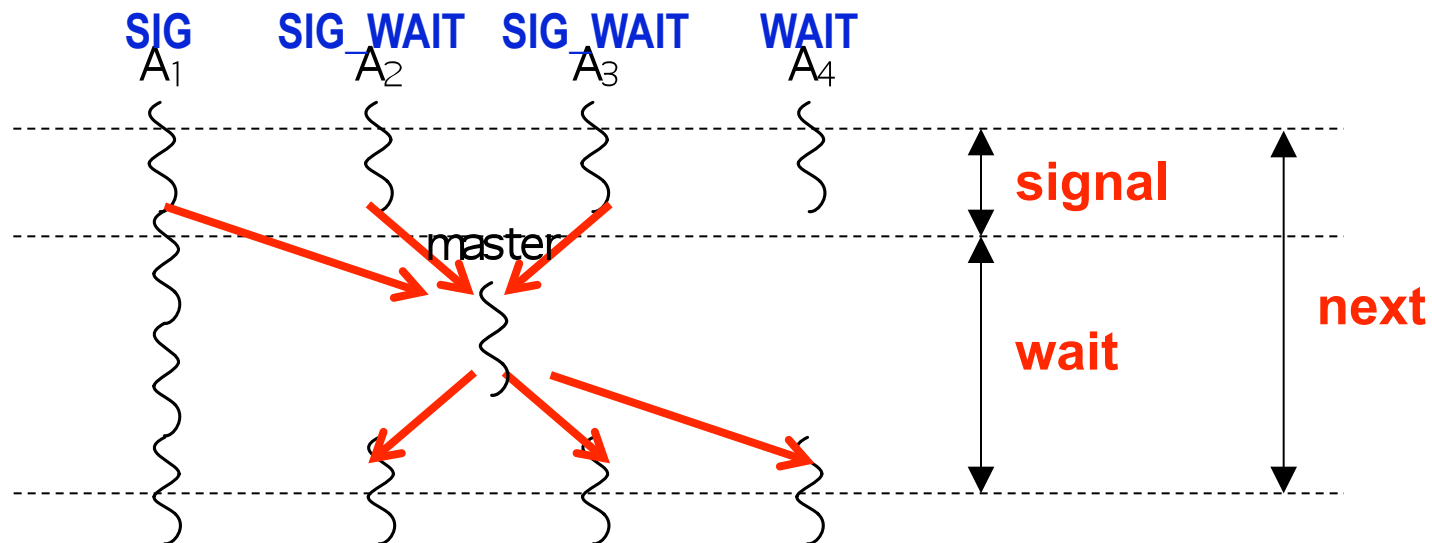
- **signal**: Non-blocking operation to notify “I reached the sync point”

- **wait**: Blocking operation to wait for other tasks' notification

next / signal / wait

next = $\begin{cases} \bullet \text{ Notify "I reached next"} & = \text{signal / ph.signal()} \\ \bullet \text{ Wait for others to notify} & = \text{wait / ph.wait()} \end{cases}$

- Synchronization semantics depends on mode
 - SIG_WAIT: **next = signal + wait**
 - SIG: **next = signal + no-op** (Don't wait for any task)
 - WAIT: **next = no-op + wait** (Don't signal any task)



- A master task is selected in tasks w/ wait capability
- It receives all signals and broadcasts a barrier completion notice

Accumulators

- **Constructs for reduction combined with phaser barrier**
- **Allocation (constructor)**
 - `accumulator(Phaser ph, accumulator.Operation op, Class type);`
 - `ph`: Host phaser upon which the accumulator will rest
 - `op`: Reduction operation
 - `sum, product, min, max, any`
 - `type`: Data type
 - `byte, short, int, long, float, double, Object (only for any)`
- **Send a data to accumulator in current phase**
 - `void put(Number data);`
- **Retrieve the reduction result from previous phase**
 - `Number get();`
 - Eager vs. lazy accumulation implementations

Phaser Accumulators for Reduction

```
phaser ph = new phaser(SIG_WAIT);  
accumulator a = new accumulator(ph, accumulator.SUM, int.class);  
accumulator b = new accumulator(ph, accumulator.MIN, double.class);
```

```
// foreach creates one task per iteration
```

```
foreach (point [i] : [0:n-1]) phased (ph<SIG_WAIT>) {
```

```
    int iv = 2*i + j;
```

```
    double dv = -1.5*i + j;
```

Must be SIG_WAIT / SIG_WAIT_SINGLE

```
    a.put(iv);
```

```
    b.put(dv);
```

Send a value to accumulator

```
    next;
```

Barrier to advance the phase

```
    int sum = a.get().intValue();
```

```
    double min = b.get().doubleValue();
```

```
    ...
```

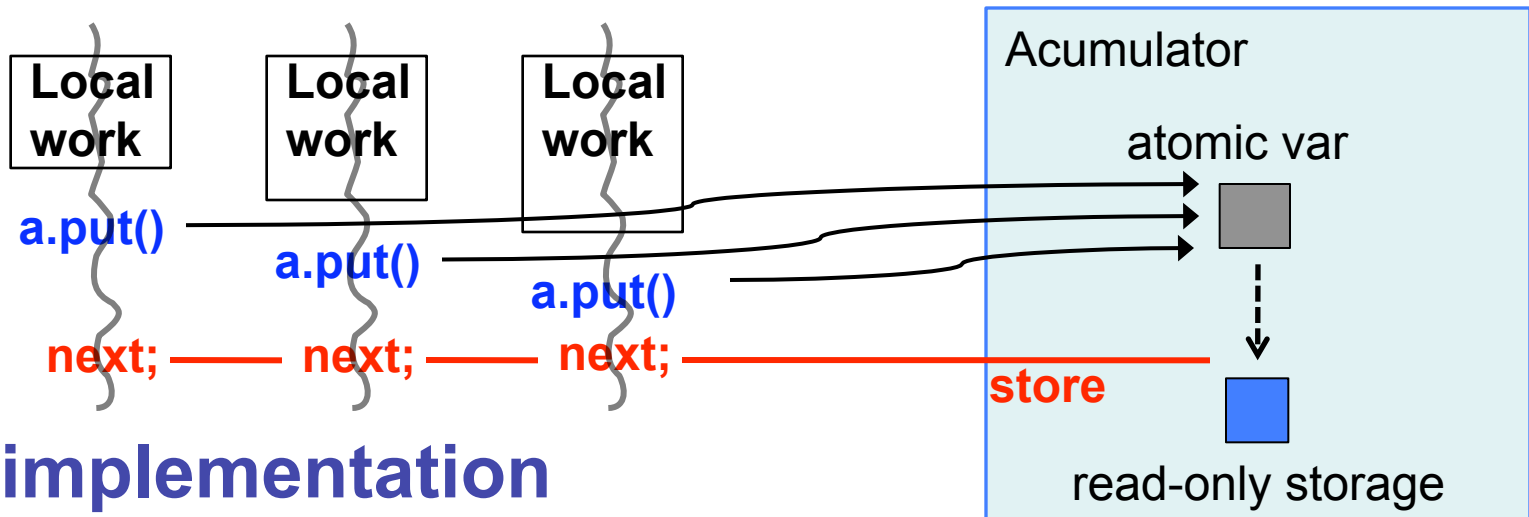
```
}
```

Get the result from *previous* phase (no race condition)

Different implementations for Accumulation

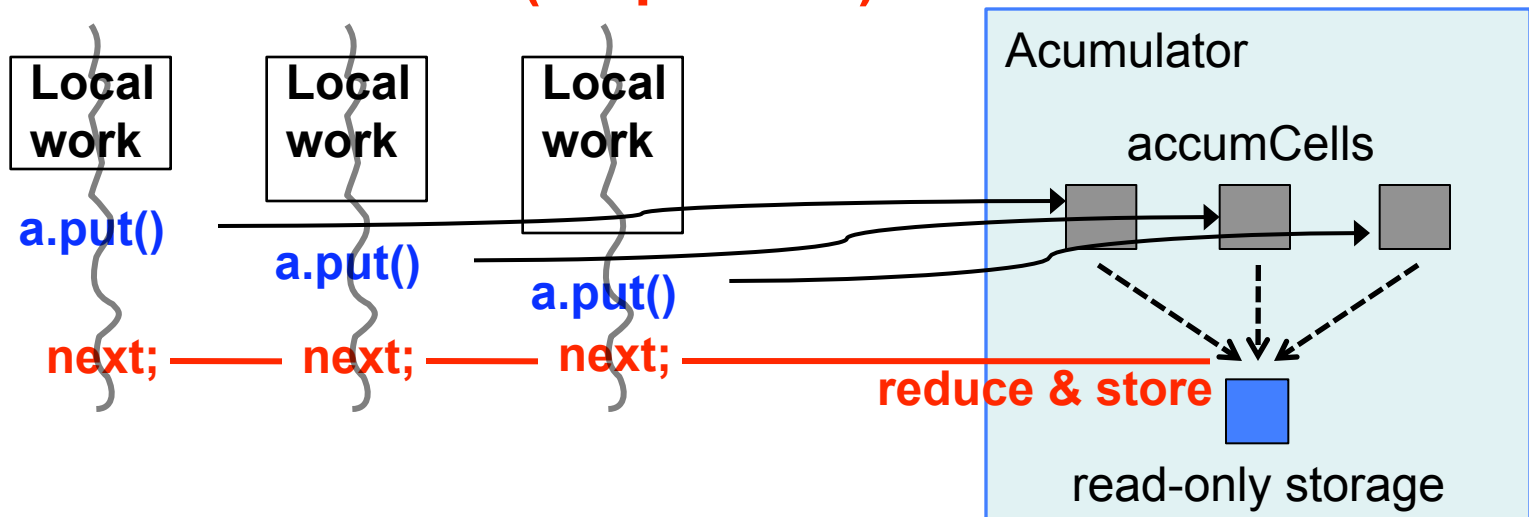
- **Eager implementation**

- Accumulation at **send (concurrent)**



- **Lazy implementation**

- Accumulation at **next (sequential)**



Outline

- Introduction
- Habanero-Java parallel constructs
 - async, finish, phasers and accumulator
- **Extensions for streaming with dynamic parallelism**
 - Phaser beams
 - Expressed streaming patterns
- **Adaptive batch optimization**
 - Runtime cycle detection
 - Adaptive batching to avoid deadlock
- **Experimental results**
- **Conclusions**

Streaming Communications

- **Producer tasks**

- Put data on stream
- Should go ahead of consumers
- Tasks on phaser in SIG mode

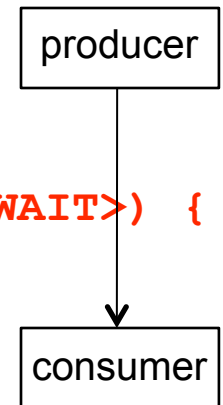
- **Consumer tasks**

- Get data from stream
- Must wait for producers
- Tasks on phaser in WAIT mode

- **Streams**

- Manage communication among tasks
 - Keep data from producers until consumers are done
 - Limited size buffer to keep data
- Accumulator to implement stream
 - Lock-step execution
 - Keep only a single data element
 - Tasks must be in SIG_WAIT

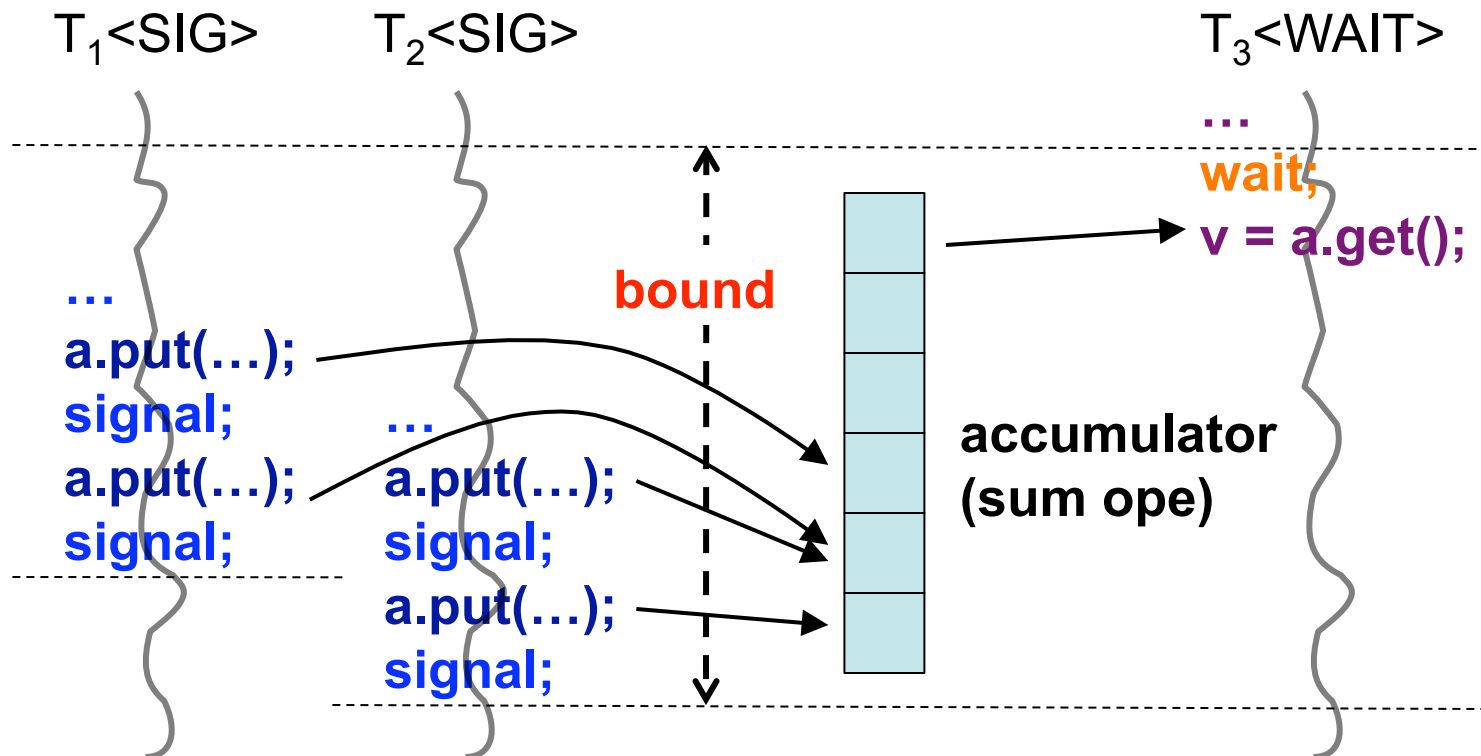
```
phaser ph = new phaser();
async phased (ph<SIG>) {
    while(...) {
        ...
        next;
        ...
    } }
async phased (ph<WAIT>) {
    while(...) {
        ...
        next;
        ...
    } }
```



Bounded Phaser Extensions

```
phaser ph      = new phaser(SIG_WAIT, bound);  
accumulator a = new accumulator(ph, SUM, double.class);
```

- **Internal buffer to accumulator**
 - Keep multiple results from bounded number of previous phases
- **Bound constraint**
 - $\# \text{ wait ops} \leq \# \text{ signal ops} \leq \# \text{ wait ops} + \text{bound size}$



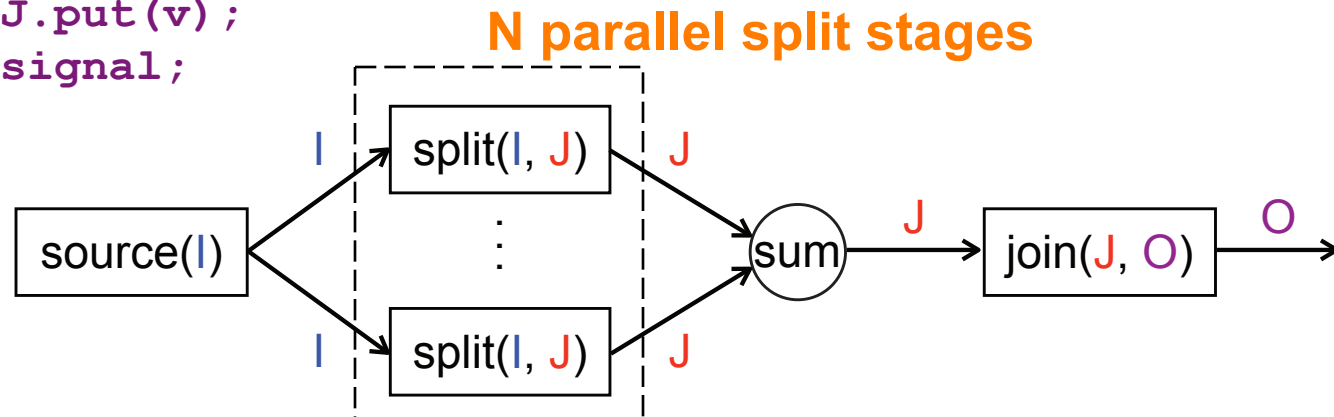
Streaming Patterns: Pipeline

```
void Pipeline() {  
    phaser phI      = new phaser(SIG_WAIT, bnd);  
    accumulator I   = new accumulator(phI, accumulator.ANY);  
    phaser phM      = new phaser(SIG_WAIT, bnd);  
    accumulator M   = new accumulator(phM, accumulator.ANY);  
    phaser phO      = new phaser(SIG_WAIT, bnd);  
    accumulator O   = new accumulator(phO, accumulator.ANY);  
    async phased (phI<SIG>)          source(I);  
    async phased (phI<WAIT>, phM<SIG>) avg(I,M);  
    async phased (phM<WAIT>, phO<SIG>) abs(M,O);  
    async phased (phO<WAIT>)          sink(O);  
}  
  
void avg(accumulator I, accumulator M) {  
    while(...) {  
        wait; wait;          // wait for two elements on I  
        v1 = I.get(0);        // read first element  
        v2 = I.get(-1);       // read second element (offset = -1)  
        M.put((v1+v2)/2);     // put result on M  
        signal;  
    }  
}
```



Streaming Patterns: Split-join

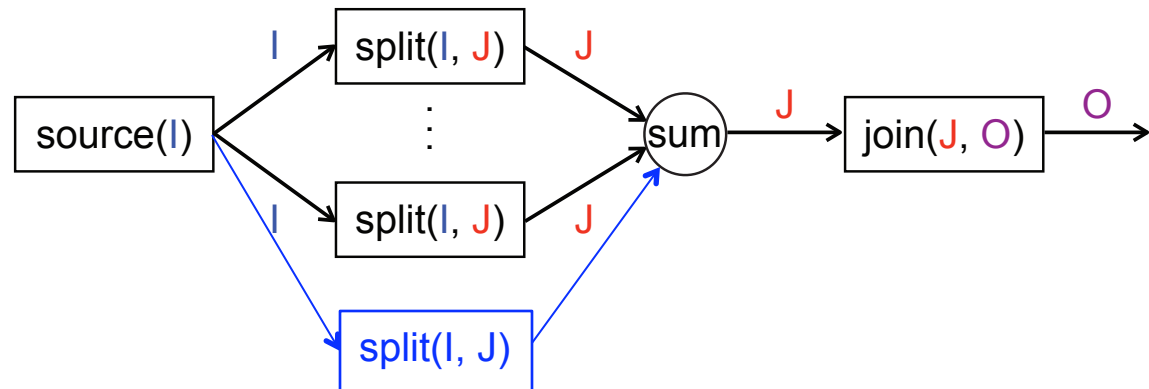
```
void Splitjoin() {  
    phaser phI          = new phaser(SIG_WAIT, bnd);  
    accumulator I       = new accumulator(phI, accumulator.ANY);  
    phaser phJ          = new phaser(SIG_WAIT, bnd);  
    accumulator J       = new accumulator(phJ, accumulator.SUM);  
  
    async phased (phI<SIG>)          source(I);  
    foreach (point [s] : [0:N-1])  
        phased (phI<WAIT>, phJ<SIG>) split(I, J);  
    async phased (phJ<WAIT>)          join(J);  
}  
split(I, J) {  
    while(...) {  
        wait;  
        v = foo(I.get());  
        J.put(v);  
        signal;  
    }  
}
```



General Streaming Graphs with Dynamic Parallelism

- Dynamic split-join

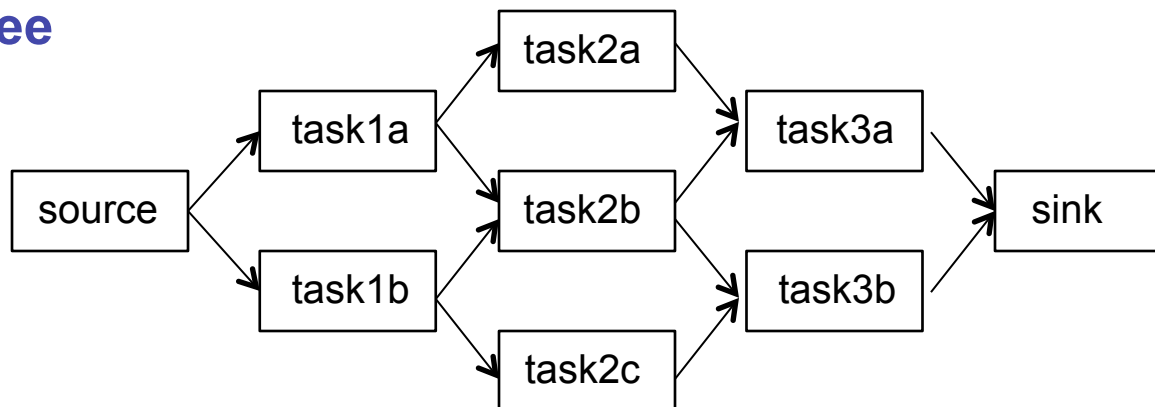
```
dynamicSplit(I, J) {  
  while(...) {  
    if (spawnNewNode()) async phased dynamicSplit(I, J);  
    if (terminate()) break;  
    wait; ...  
  }  
}
```



stages are spawned/terminated dynamically

- Dynamic pipeline

- Tree



Outline

- Introduction
- **Habanero-Java parallel constructs**
 - async, finish, phasers and accumulator
- **Extensions for streaming with dynamic parallelism**
 - Phaser beams
 - Expressed streaming patterns
- **Adaptive batch optimization**
 - Runtime cycle detection
 - Adaptive batching to avoid deadlock
- **Experimental results**
- **Conclusions**

Batch Optimization for Acyclic Graph

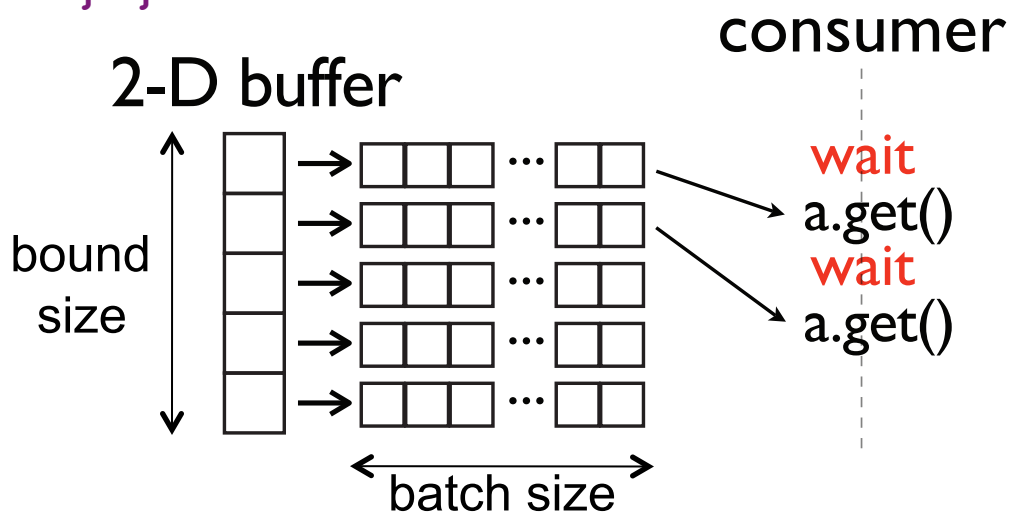
- Reduce communication overhead by factor of batch size
- Deadlock due to producer-consumer cycle

// Non-batched code

```
async phased
(ph1<WAIT>, ph2<SIG>) {
  while(...) {
    wait;
    v = foo(a1.get());
    a2.put(v);
    signal;
  }
}
```

// Batched code

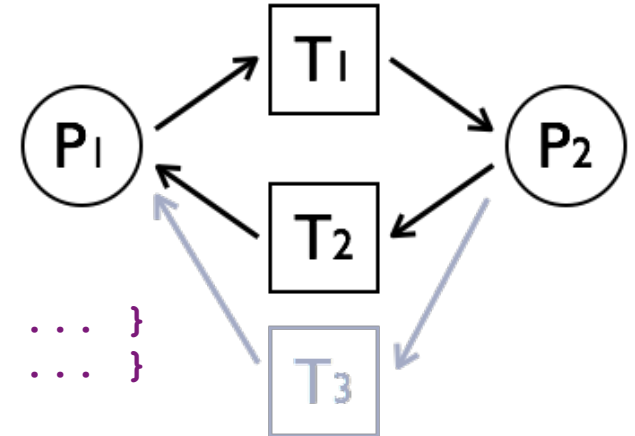
```
async phased
(ph1<WAIT>, ph2<SIG>) {
  while(...) {
    if (batch1.empty()) {
      wait;
      batch1 = a1.get();
    }
    v = foo(batch1.pop());
    batch2.push(v);
    if (batch2.full()) {
      a2.put(batch2);
      signal;
    }
  }
}
```



Adaptive Batch Optimization

- **Simple cycle example**

```
finish {  
    // Parent (root) task create phasers  
    phaser P1 = new phaser(SIG_WAIT);  
    phaser P2 = new phaser(SIG_WAIT);  
    async phased (P1<WAIT>, P2<SIG>) { // T1 ... }  
    async phased (P2<WAIT>, P1<SIG>) { // T2 ... }  
}
```



- **Adaptive batching**

- Provide batched code and non-batched code (defined in macro)
- Runtime cycle detection
 - D. Yellin, “Speeding up dynamic transitive closure for bounded degree graphs” , Acta Informatica, 30:369–384, 1993
- Switch to non-batched code if cycle is detected

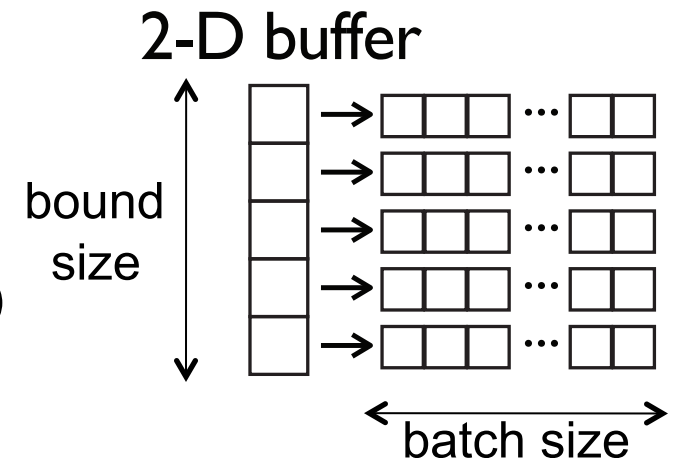
- **Capability rule in registration mode**

- Child task’s mode must be subset of parent task’s mode
 - Child task doesn’t introduce new cycle, trace only parent

Experimental Setup

- **Platforms**

- Intel Xeon E7330
 - 2.4GHz 16-core (4 Core-2-Quad)
- Sun UltraSPARC T2
 - 1.2GHz 64-thread (8-core x 8-thread/core)
- IBM Power7
 - 3.55GHz 32-core (SMT turned off)



- **Experimental variants**

- MIT StreamIt compiler & runtime 2.1.1
 - C-based implementation
 - Always apply batch optimization (assumes acyclic stream graph)
 - Batch size = 10,000, bound = unlimited (std::queue)
- Habanero-Java phasers
 - Java-based implementation
 - Adaptive batching (no constraint on stream graph structure)
 - Batch size = 10,000, bound = 8

Experimental Setup

- **Microbenchmarks**

- Push/pop microbenchmark
 - Single-producer / single-consumer
 - Throughput of streaming communication
- Thread-ring (the Computer Language Benchmarks Game)
 - Threads are linked in a ring (cycle structure)
 - A token is passed around
 - Efficiency of runtime cycle detection

- **Application benchmarks**

- Filterbank, FMRadio, BeamFormer (StreamIt benchmarks)
 - Acyclic graph structure
 - Static stream graph w/o dynamic parallelism
- Sieve of Eratosthenes
 - Find prime numbers from input stream (increasing integers)
 - Dynamic pipeline / dynamic split-join

Microbenchmarking Results

- **Push/pop: 1-producer / 1-consumer**

- # operations per second
- Busywait-based phaser vs. lock-based StreamIt

	Xeon	T2	Power7
StreamIt (batch)	114.0×10^6	21.7×10^6	33.1×10^6
Phaser (non-batch)	11.0×10^6	2.7×10^6	8.4×10^6
Phaser (adaptive batch)	148.2×10^6	24.5×10^6	299.4×10^6

- **Thread-ring: Cyclic streaming graph**

- Average time per hop [microseconds]
- Small overhead for adaptive batching

	Xeon	T2	Power7
Java original	9.4 μ s	16.3 μ s	11.9 μ s
StreamIt (batch)	N/A	N/A	N/A
Phaser (non-batch)	2.2 μ s	2.7 μ s	2.9 μ s
Phaser (adaptive batch)	2.2 μ s	2.7 μ s	3.0 μ s

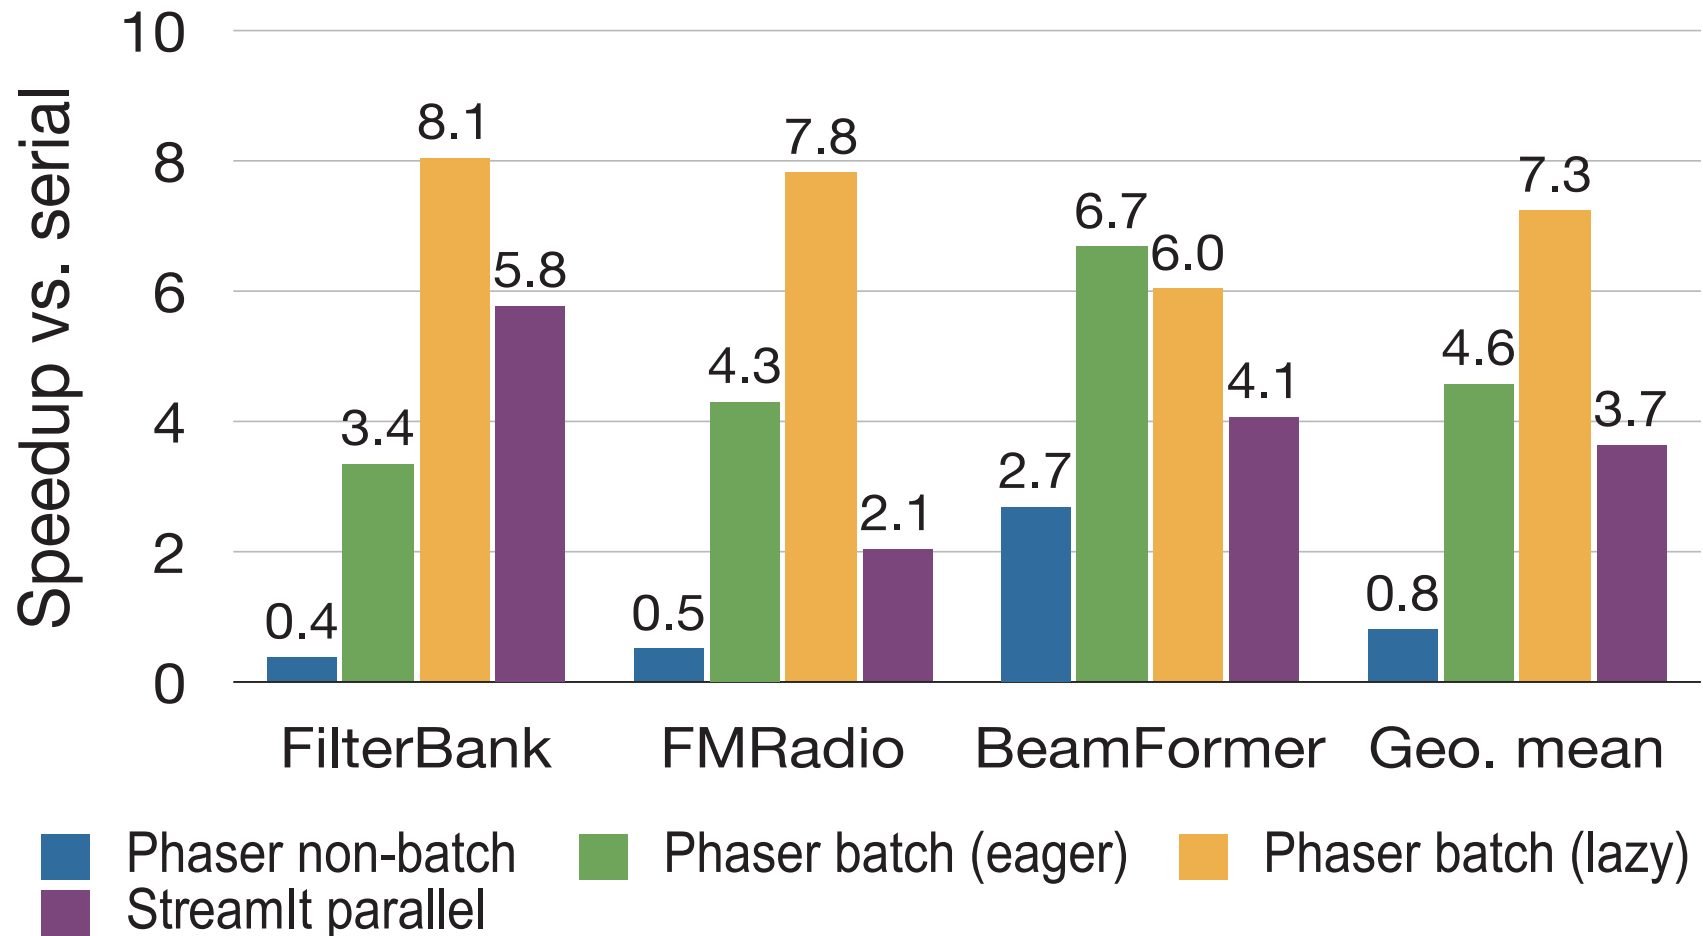
Summary for StreamIt Benchmarks

Benchmark	variant	Xeon	T2	Power7
FilterBank	Java serial	11.4 sec	175.6 sec	15.1 sec
	HJ parallel (phaser)	1.4 sec	23.9 sec	3.4 sec
	StreamIt serial	8.9 sec	41.2 sec	1.9 sec
	StreamIt parallel	1.5 sec	6.7 sec	5.4 sec
FMRadio	Java serial	25.3 sec	288.1 sec	26.6 sec
	HJ parallel (phaser)	3.2 sec	20.7 sec	4.8 sec
	StreamIt serial	7.6 sec	470.3 sec	5.9 sec
	StreamIt parallel	3.7 sec	21.2 sec	8.0 sec
BeamFormer	Java serial	19.1 sec	258.7 sec	20.7 sec
	HJ parallel (phaser)	3.2 sec	35.2 sec	6.0 sec
	StreamIt serial	6.4 sec	86.8 sec	8.9 sec
	StreamIt parallel	1.6 sec	13.4 sec	3.5 sec
Geo-mean (speedup vs. Java serial)	HJ parallel (phaser)	7.3×	9.1×	4.4×
	StreamIt serial	2.3×	2.0×	4.4×
	StreamIt parallel	8.5×	19.0×	3.8×

- HJ parallel: Lazy implementation policy for accumulator
- StreamIt serial (C-based): 2.0x – 4.4x faster Java serial

Scalability (vs. each sequential base lang.)

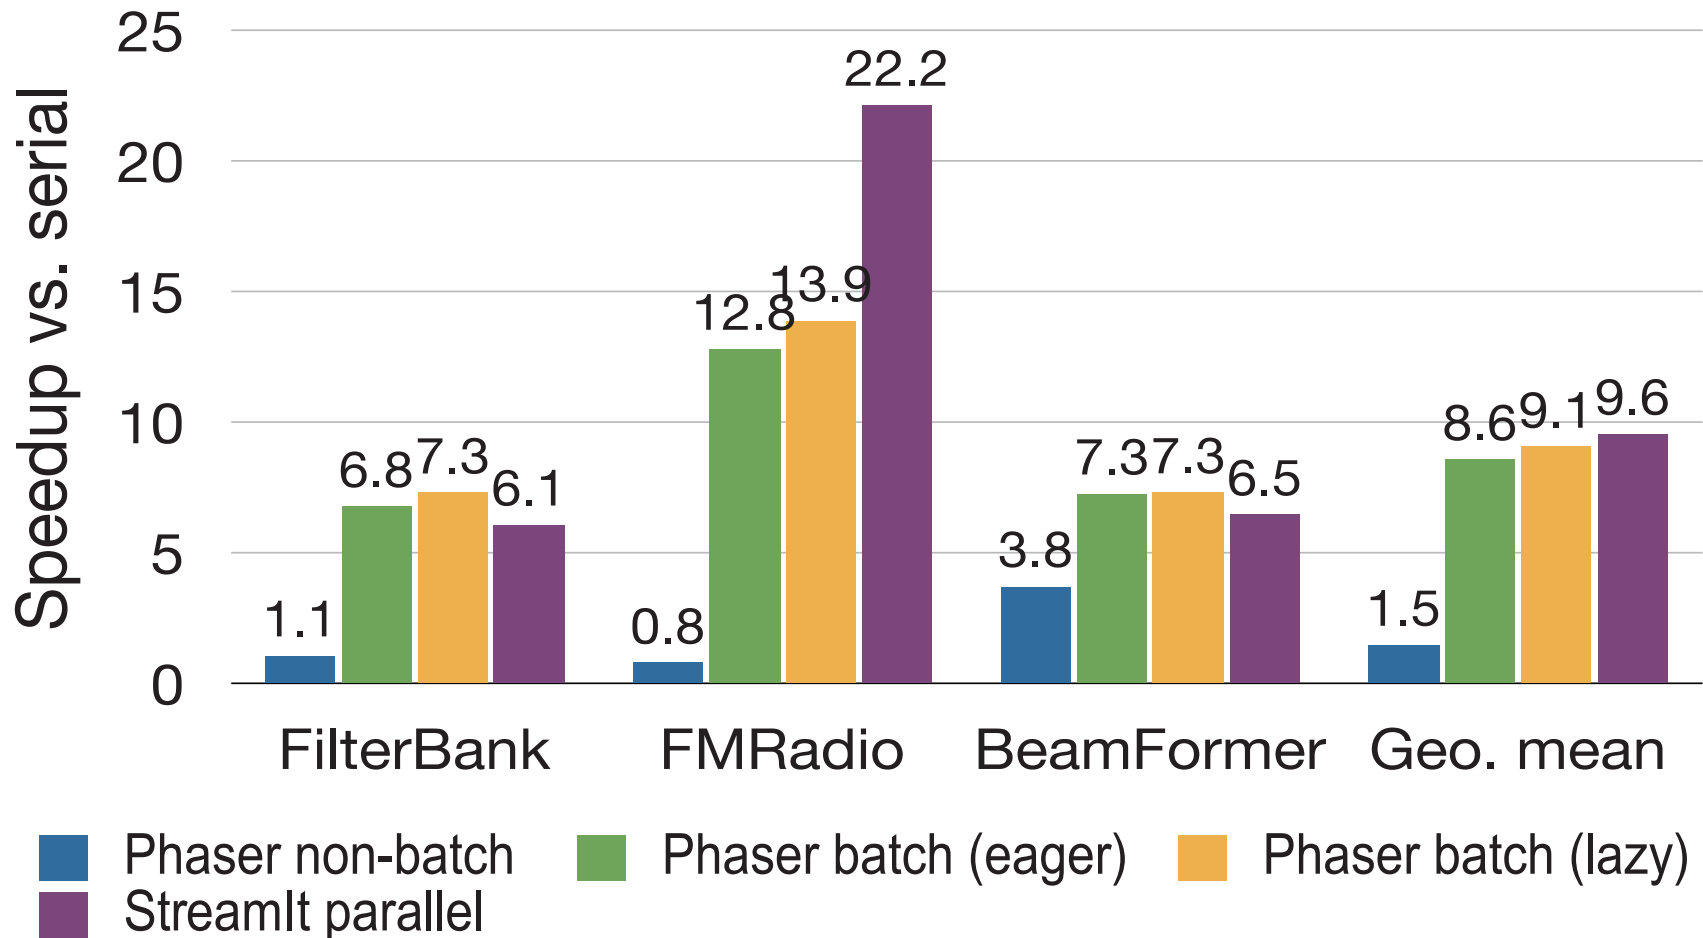
2.4GHz 16-core Intel Xeon



- Better scalability due to synchronization efficiency of phasers
- Accumulator implementation: Lazy policy > Eager policy

Scalability (vs. each sequential base lang.)

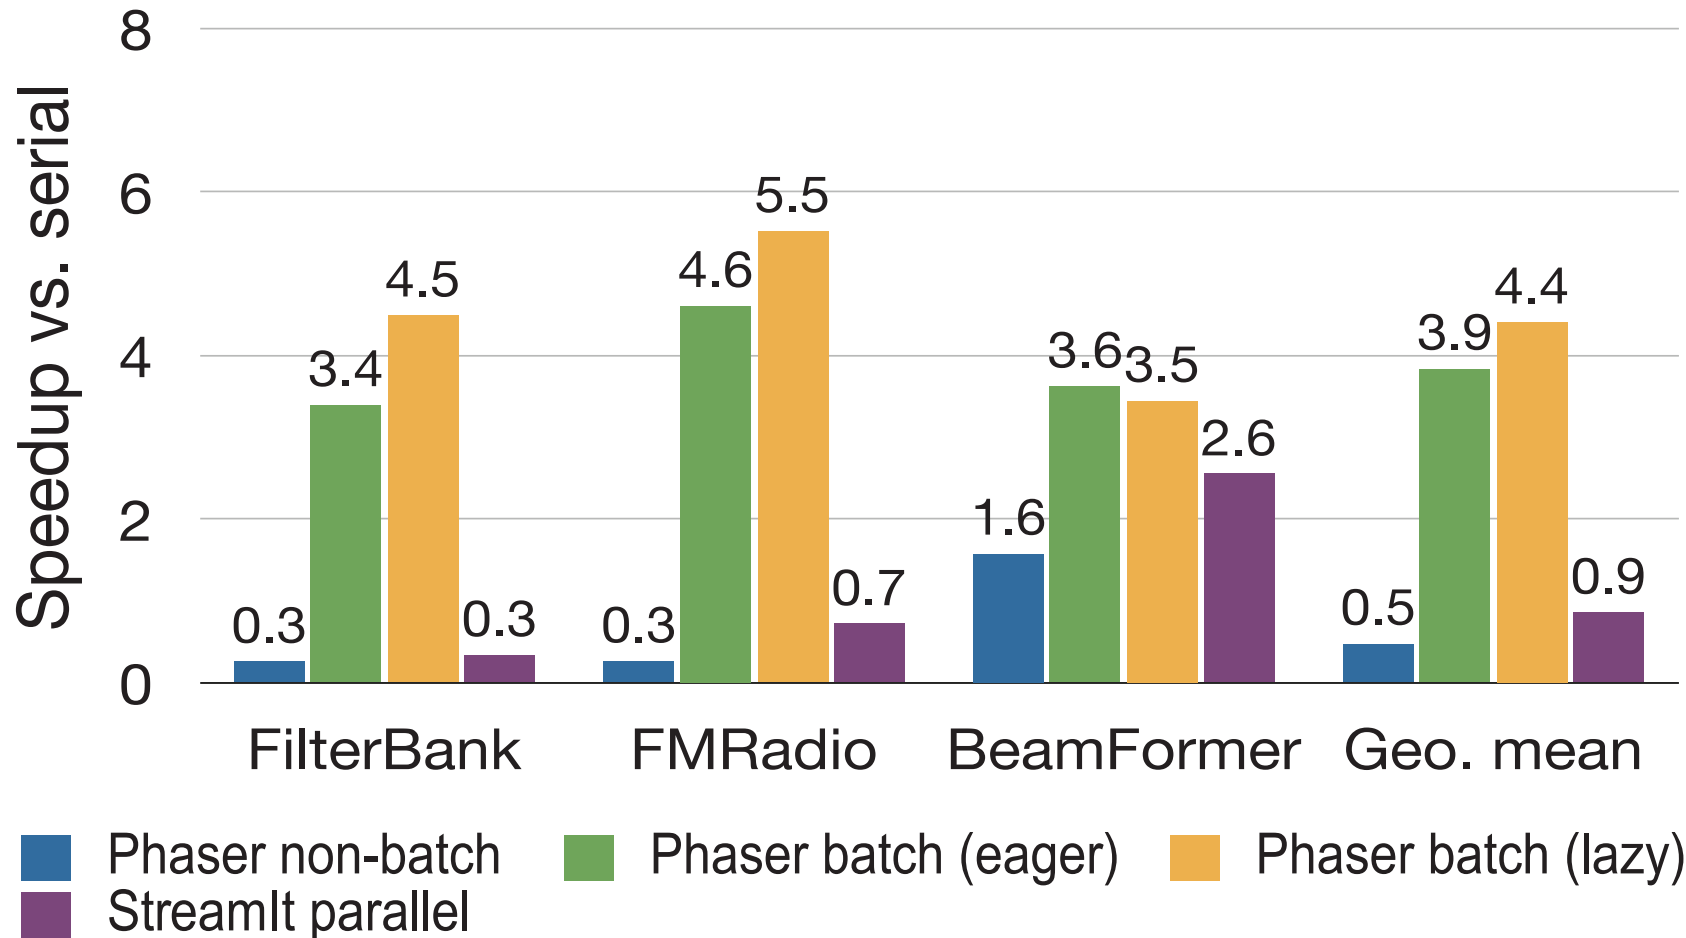
1.2GHz 8-core x 8-thread/core Sun T2



- Scalability of StreamIt is better than phasers
- Accumulator implementation: Lazy policy \approx Eager policy

Scalability (vs. each sequential base lang.)

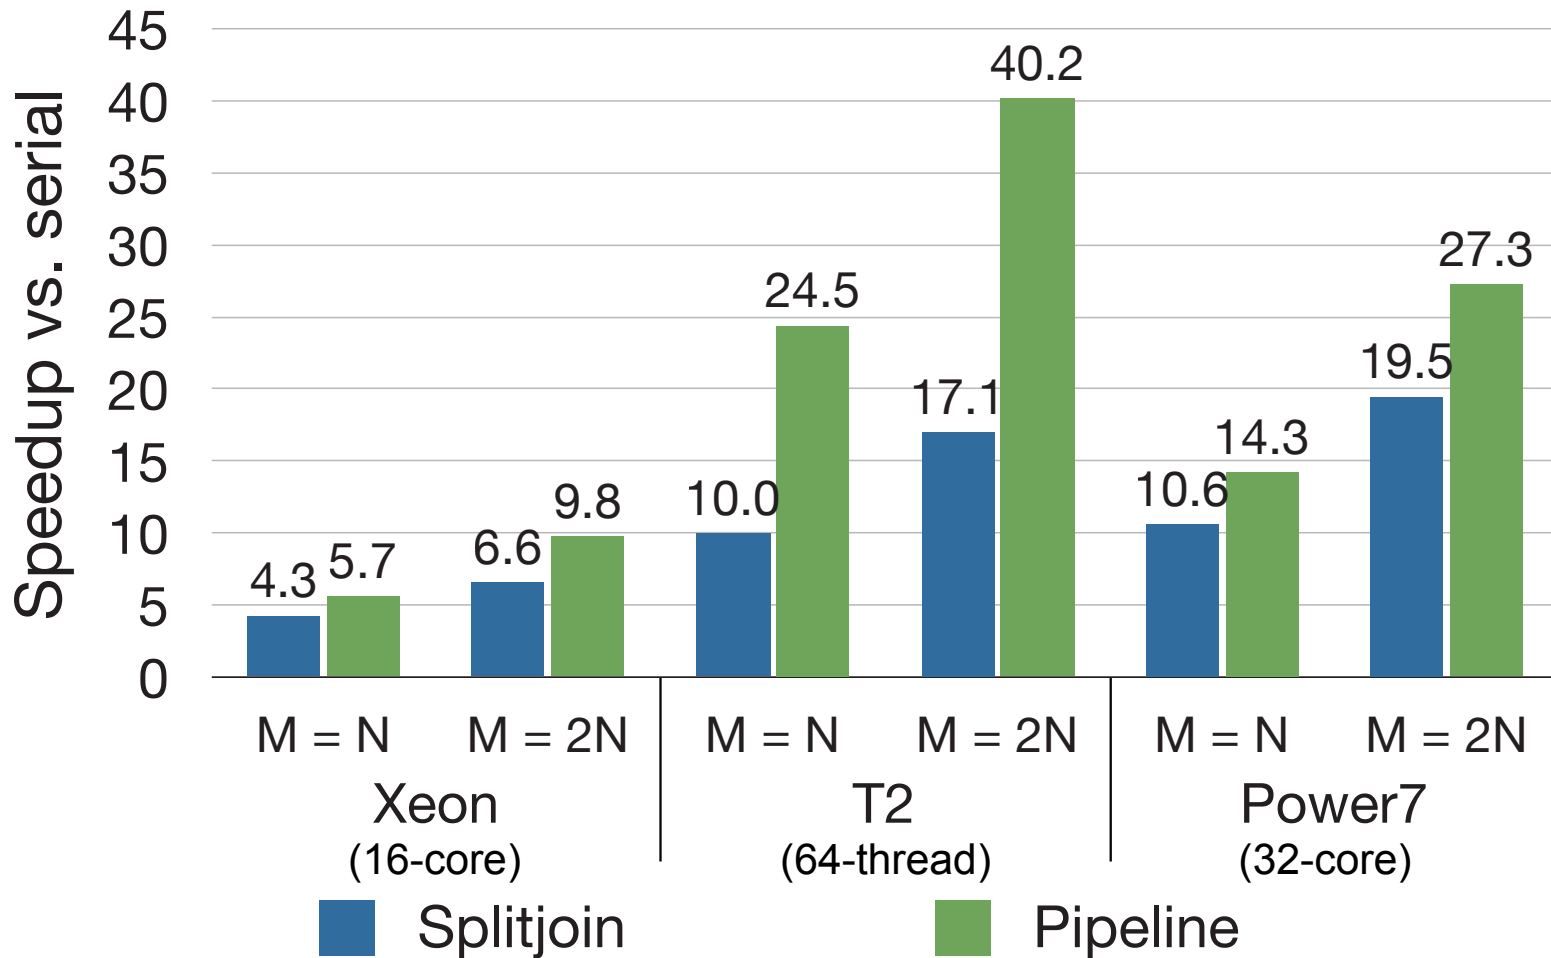
3.55GHz 32-core IBM Power7



- Better scalability due to synchronization efficiency of phasers
- Accumulator implementation: Lazy policy > Eager policy

Sieve of Eratosthenes

(Integration of Dynamic Task and Stream Parallelism)

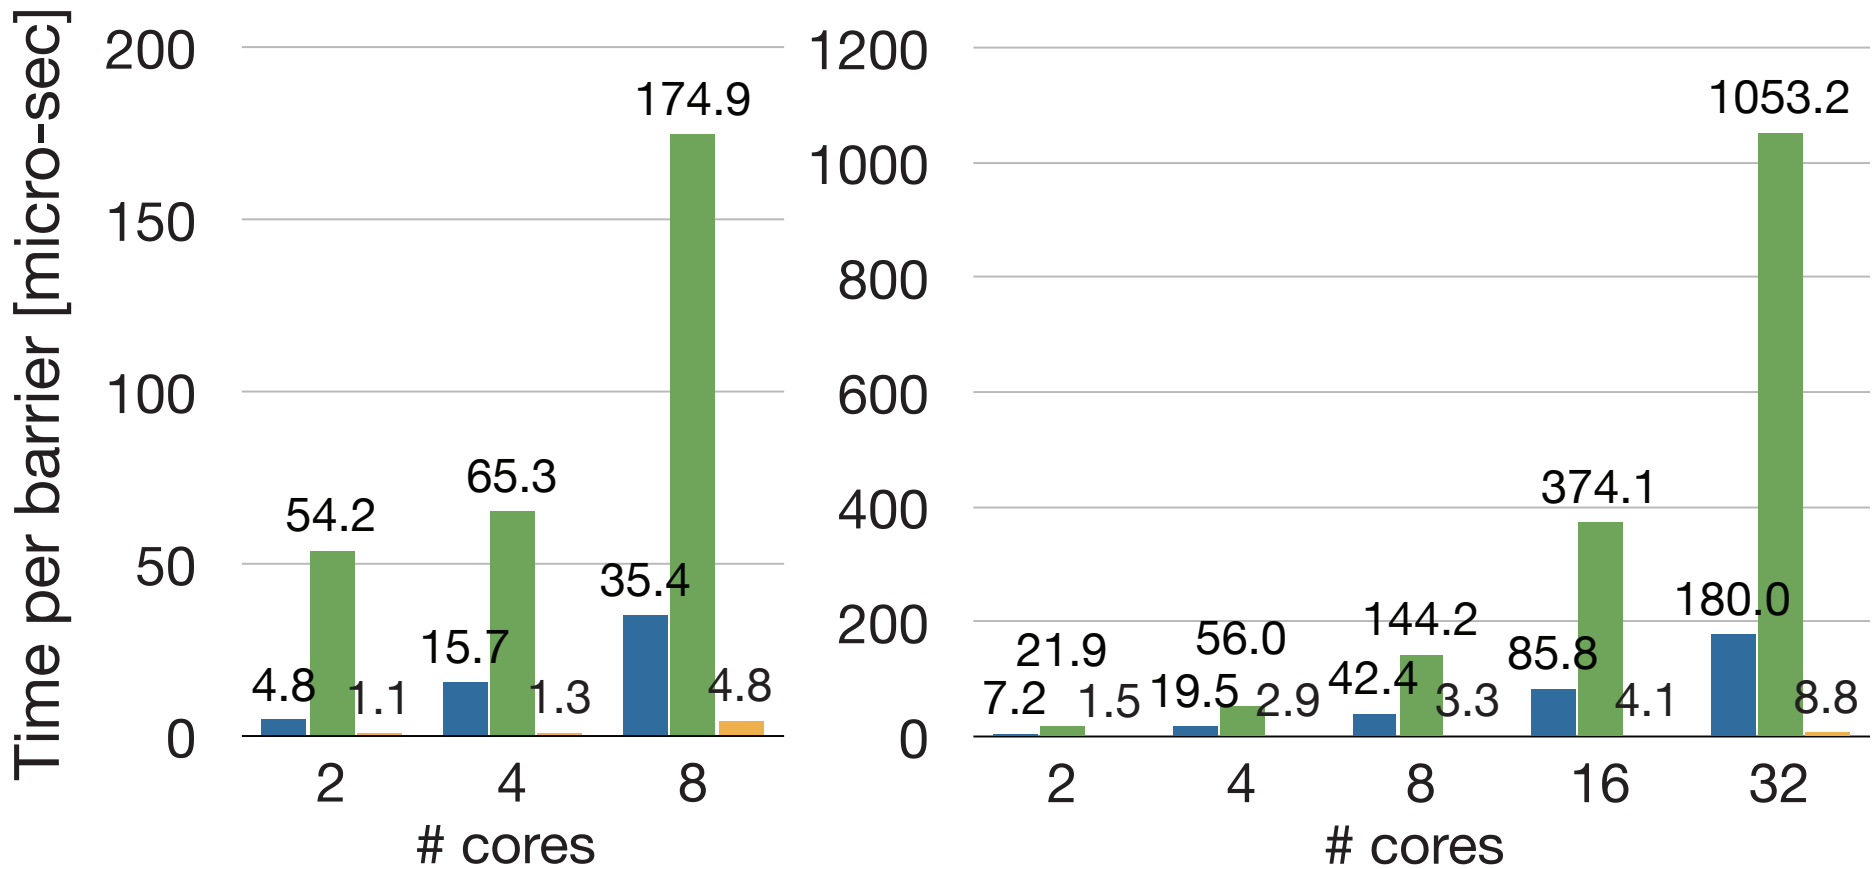


- M : Upper bound of integer in input stream
- N : Upper bound of prime number

Conclusion

- **Phaser beams for streaming computation**
 - Integrating task and stream parallelism in a programming model
 - Adaptive batching with cycle detection
- **Experimental results on three platforms**
 - Push/pop microbenchmark (vs. C-based StreamIt)
 - 1.3x faster on Xeon, 1.1x on T2 and 9.0x on Power7
 - StreamIt benchmarks (vs. each sequential base lang.)
 - HJ phasers: 7.3x on Xeon, 9.1x on T2, and 4.4x faster on Power7
 - StreamIt: 3.7x on Xeon, 9.6x on T2, and 0.9x on Power7
 - Sieve of Eratosthenes (vs. sequential Java)
 - Up to 9.8x on Xeon, up to 40.2x on T2, and up to 27.3x on Power7
- **Future work**
 - Dynamic selection of eager or lazy policy
 - Static compiler optimizations, e.g., batch code generation and graph partitioning
 - Support of phaser functionality in X10 programming language

Barrier Performance of CyclicBarrier, Clocks and Phasers



(a) Nehalem

(b) Power7

■ JUC CyclicBarrier ■ X10 v2.1 clock (c++) ■ Phasers

- Nehalem: Intel Corei7 2.4GHz 2 quad-core processor
- Power7: IBM Power7 3.55GHz