

Parallel Programming: Design of an Overview Class

Christoph von Praun
University of Applied Sciences
Nuremberg, Germany
praun@acm.org

This work was supported by an IBM Innovation Award grant.


Summary

- Design of a 3rd year introductory parallel programming class in the Bachelor curriculum: 'Orientation' class
- Key characteristics of the class
 - Organization of topics follows the **Tiers of parallelism**
 - Uses programming language **X10**
 - Strong focus on **lab sessions**
- Teaching materials are available online

Computer Science Master + Bachelor Curriculum


Elective classes

Semester



The Art of
Multiprocessor
Programming

Graphics Programming
with CUDA

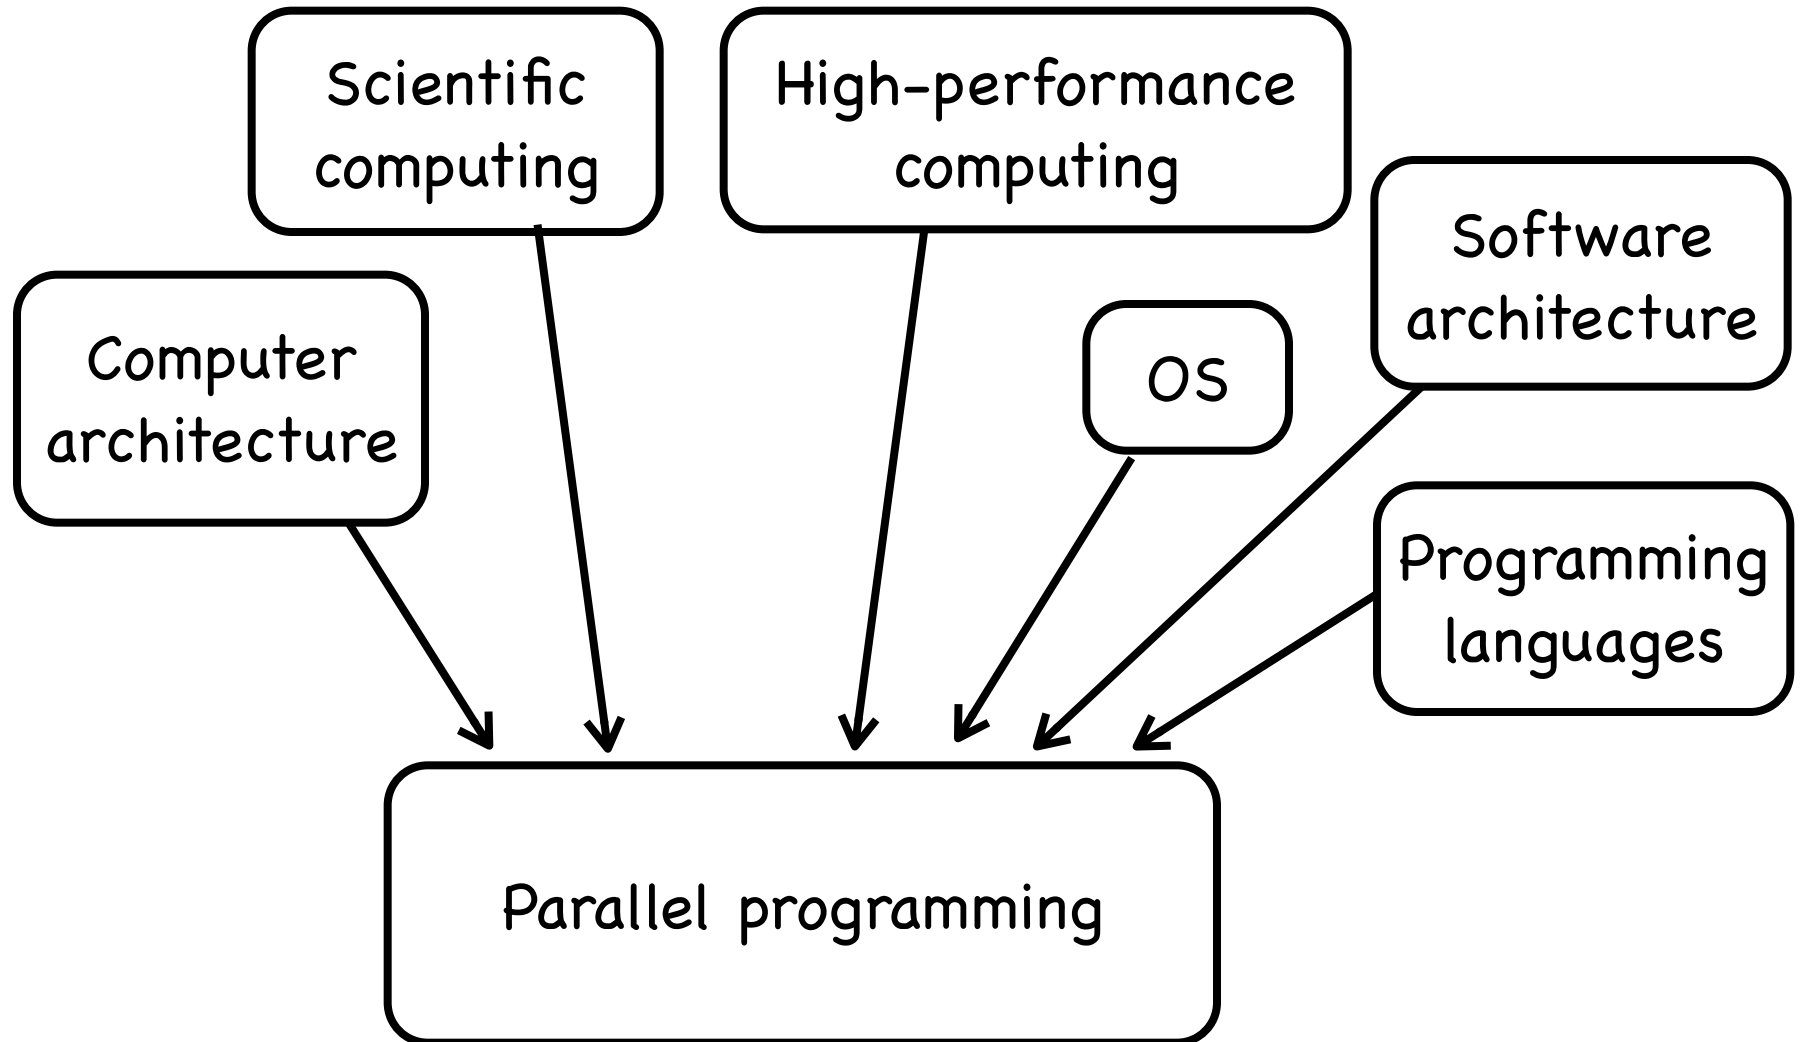


Scientific Computing

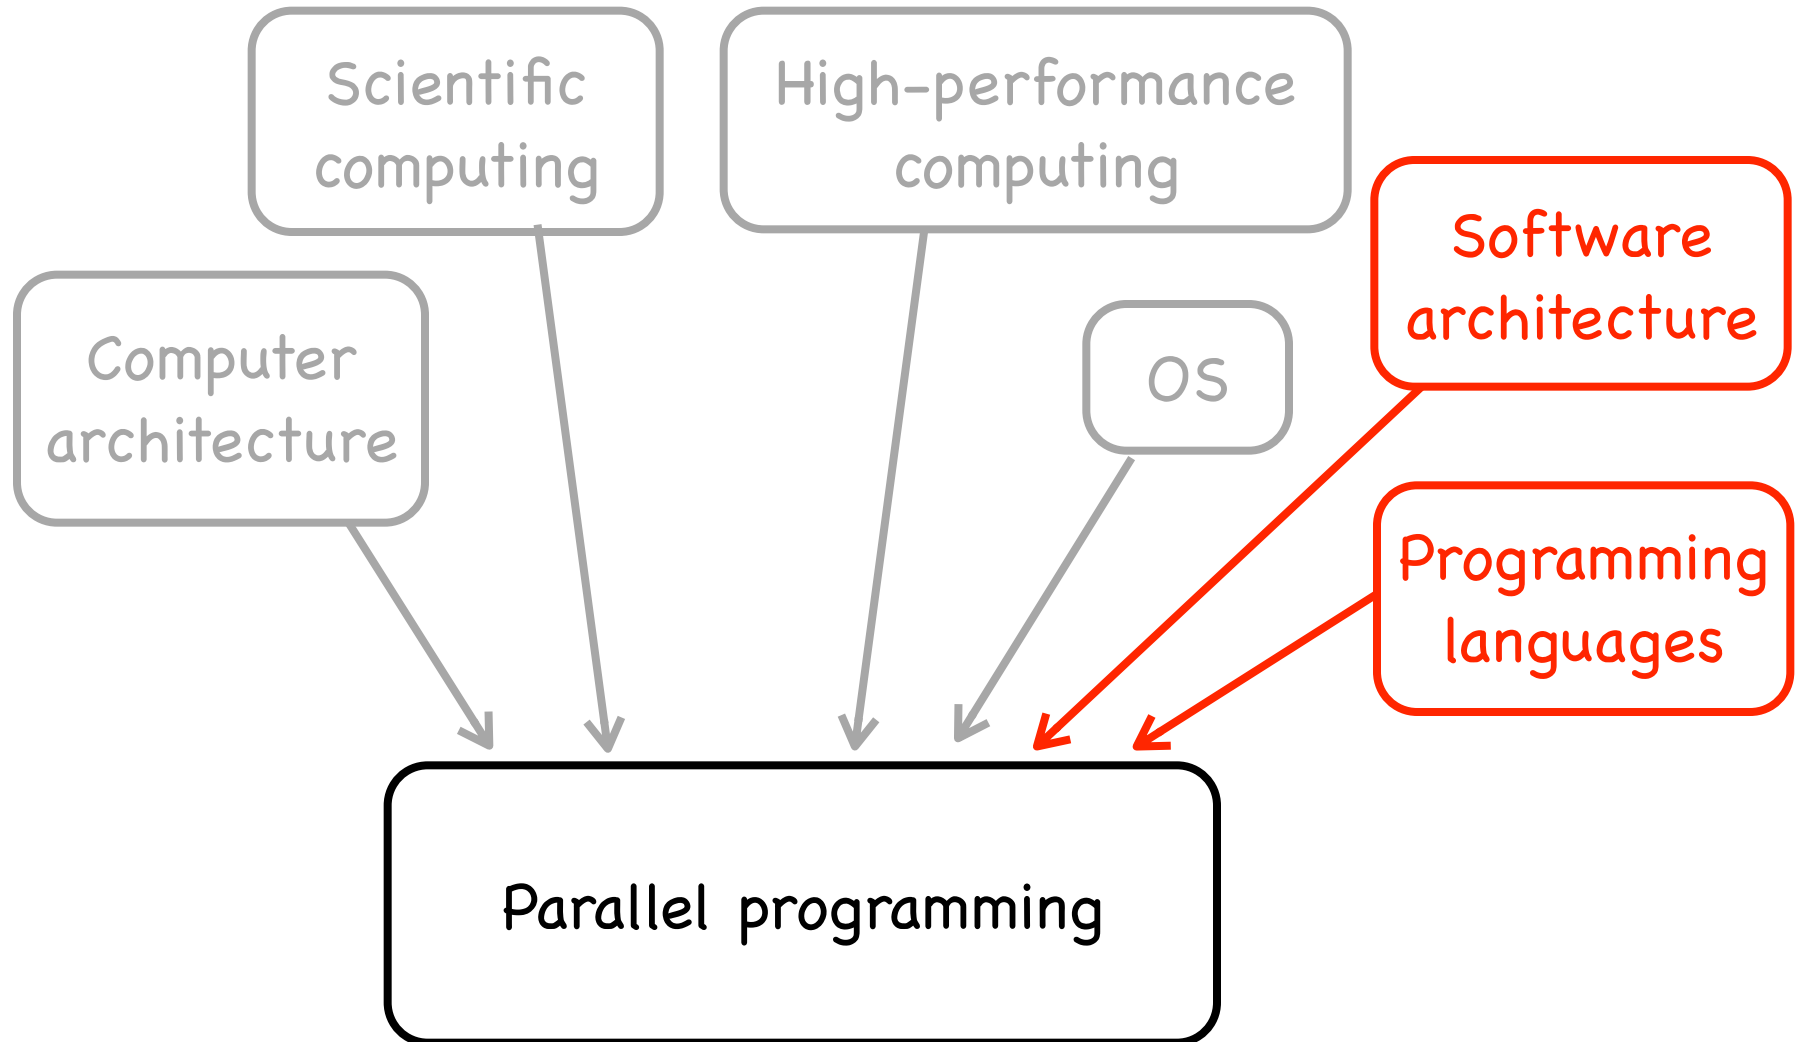
Parallel Programming

'Orientation' class

Influences on parallel programming classes



Influences on parallel programming classes



Outline

- **Tiers of parallelism**
- Course structure and contents
- Role of X10
- Student feedback and experience

Tiers of parallelism

- Original idea due to Michael L. Scott:
 - “Don’t start with Dekker’s algorithm ...” [1]
 - “Making the simple case simple” [2]
- Development of parallel software (parallel programming) can be based on techniques at different abstraction layers
 - progressively less complexity at higher abstraction layers

parallelization

techniques

(1) automatic or implicit

parallelizing compiler

(2) deterministic

fully independent computations or serialization


(3) explicitly synchronized (data race free)

critical sections, transactions

(4) low-level (with race conditions)

implementation of threads, synchronization mechanisms, non-blocking data structures

high-level
(simpler)



low-level
(more complex)

parallelization

techniques

(1) automatic or implicit

parallelizing compiler

(2) deterministic

fully independent computations or serialization

(3) explicitly synchronized (data race free)

critical sections, transactions

(4) low-level (with race conditions)

implementation of threads, synchronization mechanisms, non-blocking data structures

parallelization

techniques

(1) automatic or
implicit

parallelizing compiler

(2) de

Goal of the class:

Students should be conscious about
'their' tier when developing a parallel
program.

tion

(3) ex
synch
(data

Encourage students to move
programming activity to higher tiers in
the abstraction hierarchy.

(4) lo
(with race
conditions)

synchronization mechanisms,
non-blocking data structures

Tier-1: automatic or implicit parallelism

- **Auto-parallelization** through compilers
- **Parallel kernels:** parallelism encapsulated in libraries
 - LAPACK, etc.
- **Parallel frameworks:** framework organizes parallelism, synchronization and communication, programmer supplies sequential kernels
 - Map-reduce
 - Web application frameworks, e.g. WebSphere
 - etc.

Tier-1: automatic or implicit parallelism

- **Auto-parallelization** through compilers
- **Parallel kernels:** parallelism encapsulated in libraries
 - LAPACK, etc.
- **Parallel frameworks:** framework organizes parallelism, synchronization and communication, programmer supplies sequential kernels
 - Map-reduce
 - Web application frameworks, e.g. WebSphere

Sequential semantics.

Tier-2: deterministic parallelism

- **Independent computations:**
 - parallel array languages (FORALL loops)
 - parallel containers (e.g., STAPL, Intel Concurrent Collections, Hierarchically Tiled Arrays)
- Concurrent computations with dependencies that follow **deterministic idioms:**
 - reduction, scan

Tier-2: deterministic parallelism

- **Independent computations:**
 - parallel array languages (FORALL loops)
 - parallel containers (e.g., STAPL, Intel Concurrent Collections, Hierarchically Tiled Arrays)
- Concurrent computations with dependencies that follow **deterministic idioms:**
 - reduction, scan

Semantics through serialization + sequential reasoning.

Tier-3: explicitly synchronized, data-race-free

Three principal programming models

- **Event-based**
- **Thread-parallel with shared memory**
 - critical sections
 - condition variables
- **Message-based**
 - send/receive
 - collective communication

Tier-3: explicitly synchronized, data-race-free

Three principal programming models

- **Event-based**
- **Thread-parallel with shared memory**
 - critical sections
 - condition variables
- **Message-based**
 - send/receive
 - collective communication

Semantics through interleaving
of program blocks

Tier-4: low-level, with race conditions

- Programming with shared memory
 - atomic load and store
 - atomic compare and swap
- Platform-specific (Java, X86, ...)
- Sequential consistency is often a simplifying assumption
 - e.g. teaching Dekker's algorithm

Tier-4: low-level, with race conditions

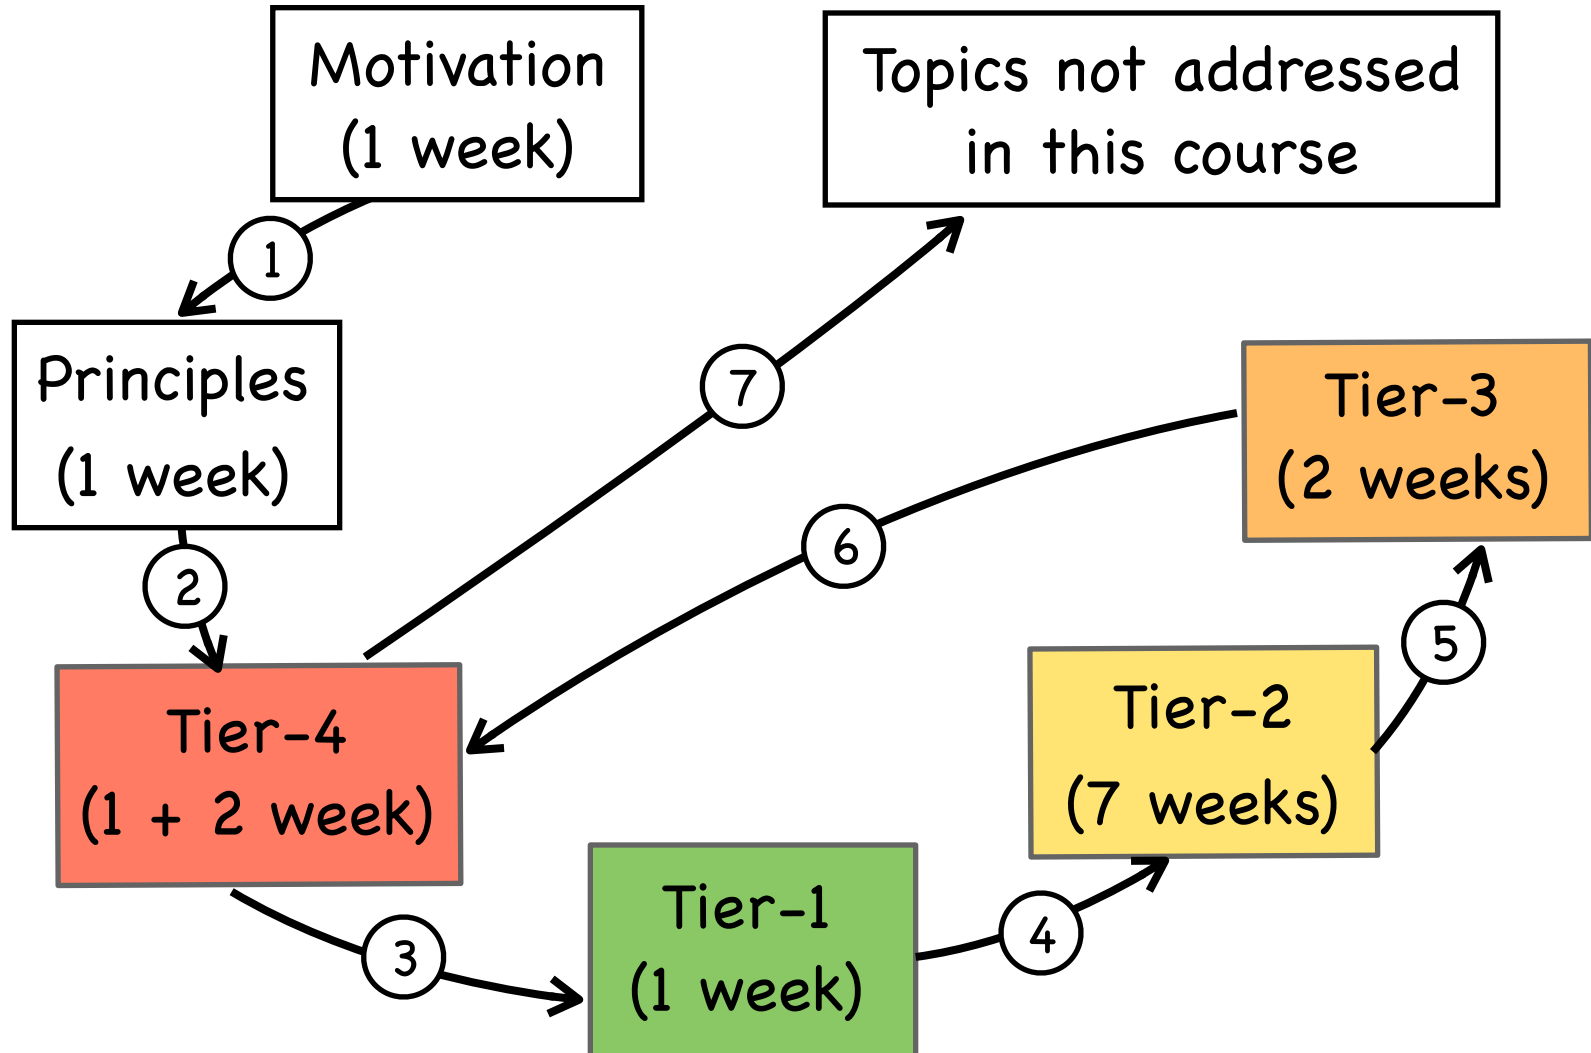
- Programming with shared memory
 - atomic load and store
 - atomic compare and swap
- Platform-specific (Java, X86, ...)
- Sequential consistency is often a simplifying assumption
 - e.g. teaching Dekker's algorithm

Semantics through interleaving
of statements, possibly not
sequentially consistent

Outline

- Tiers of parallelism
- **Course structure and contents**
- Role of X10
- Student feedback and experience

Roadmap of the class (15 weeks)



Motivation (1 week)

- Hardware trend:
 - Moore's Law continues
 - frequency scaling limited by power density: multicores
- Performance: Software need to be parallel
 - challenges (Amdahl's Law)
 - opportunities (Gustafson's Law)
- Energy: Throughput-oriented computing can save energy

Lab session

- Pencil and paper

Principles (1 week)

- Simple model for concurrent computations
 - partial orders of operations
 - synchronization vs ordinary operations
 - happens-before relation
- Explain semantics of X10 language constructs
 - async
 - finish, for-async
 - atomic

Lab session

- Parallel prime number testing

Tier-4 (1 week)

[low-level with race conditions]

- Race conditions
- Non-determinacy
 - associative non-determinism (floating point)
 - atomicity violation: lost-update problem
- “Interleaving” semantics

Lab sessions

- Numeric integration

Tier-1 (1 week)

[automatic or implicit parallelism]

- Challenges of loop parallelization
 - intro to data dependencies
 - difficulties and limitations of dependence analysis on some loop scenarios
- Parallel frameworks
 - Map-Reduce, Web-applications

Lab sessions

- Development of Map-Reduce applications (framework provided)


Tier-2 (7 weeks)

[deterministic parallelism]

Patterns for algorithmic problem decomposition:

according to T. Mattson, B. Sanders, B. Massingill,
“Patterns for parallel programming”, AW 2005.

- Data parallelism
 - geometric decomposition, recursive data
 - **data locality issues**
- Task parallelism
 - task parallel, divide and conquer
 - **task scheduling / load balancing issues**

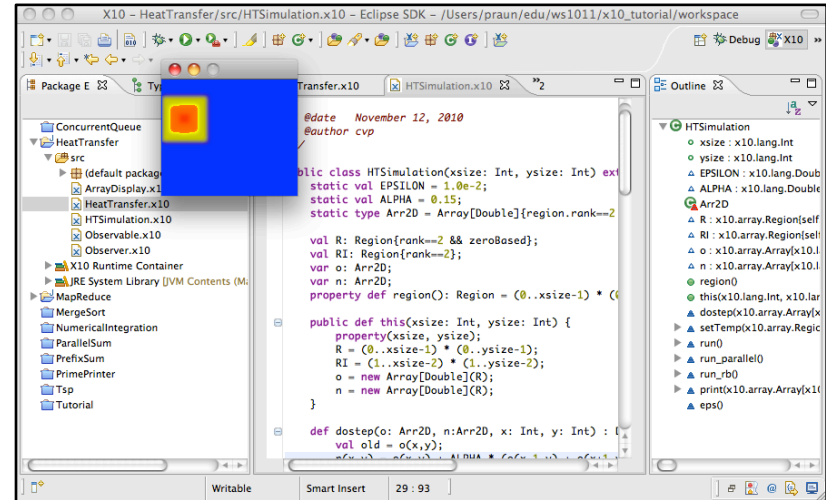


Tier-2 (7 weeks)

[deterministic parallelism]

Lab sessions

- Data parallel
 - heat-transfer
 - matrix multiply
 - algorithms for reduction and prefix-sum
- Task parallel
 - map-reduce framework implementation
 - merge-sort
 - traveling salesman



Tier-3 (2 weeks)

[explicitly synchronized]

- Pattern: Pipeline parallelism
- Producer-consumer communication through concurrent queues
 - critical sections
 - conditional synchronization

Lab sessions

- Array-based concurrent queue with explicit synchronization (atomic blocks)

Tier-4 (2 weeks)

[low-level with race conditions]

- Programming with race conditions
- Memory models (SC, TSO)

Lab sessions

- Lamport's concurrent non-blocking queue
(1 consumer / 1 produce non-blocking queue)
- Observe non-SC behavior of Java

Topics not addressed in the course


- Patterns for ...
 - ... locality / reducing data access latency
 - ... load balancing / distribution of work
 - ... enhancing parallelism
 - ... distribution data
- Performance debugging

Outline

- Tiers of parallelism
- Course structure and contents
- **Role of X10**
- Student feedback and experience

X10

- Pragmatic choice:
 - Syntax familiar to students: “extension of sequential Java”
 - Simple things can be expressed with succinct syntax
 - X10 can express programs at tiers (1)–(3)
memory model not specified -> use Java at tier (4)
- Class was not X10 ‘only’
 - students could choose their own language for projects
 - X10 language tutorial provided separately



Student feedback on X10

“The language should not be used in future classes , since parallel programming is simplified significantly, and for that reason, one does not run into issues and problems that occur when conventional programming languages are used for parallel programming.”

“Takes a while to be familiar with the type system / type inference.”


“Usability of X10 IDE needs to be improved” [March-June 2010]

Outline

- Tiers of parallelism
- Course structure and contents
- Role of X10
- **Student feedback and experience**

Student feedback (1/3)


Has the course been well-structured and did the structure support your learning?



Feedback collected from 16/21 participants.

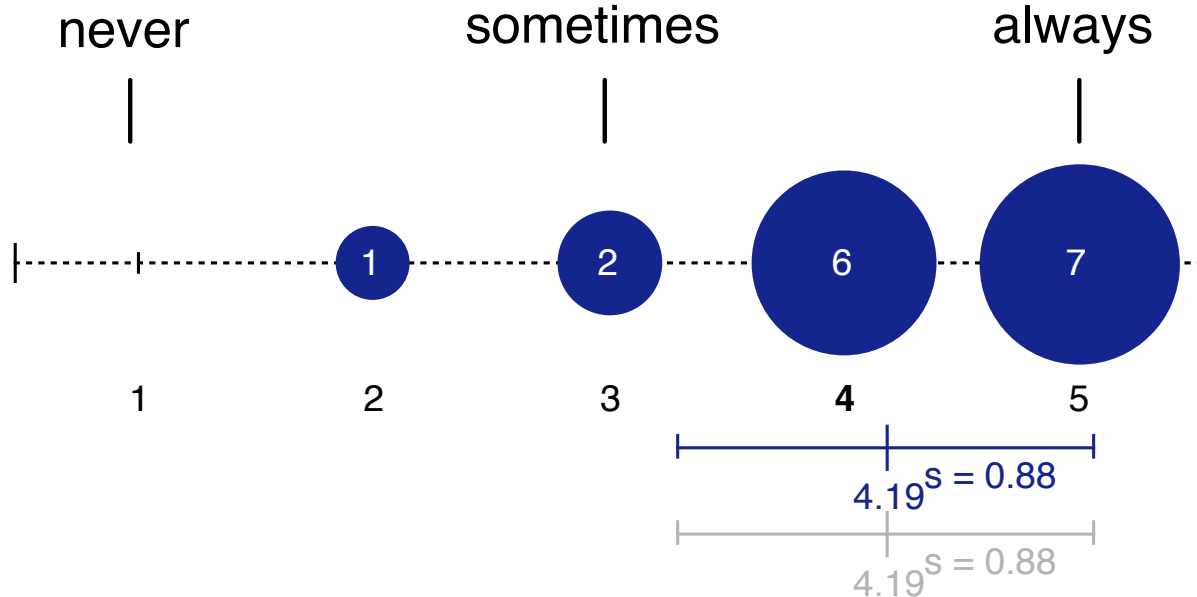
Student feedback (2/3)

The number of topics and the volume of material presented in class was ...



Student feedback (3/3)

Did the lab sessions help you to learn and understand the materials presented in class?



Criticism

- Focus of discussion on correctness, not performance
- Focus on 'higher layers' in the abstraction hierarchy
 - Less complex than lower tiers
 - Assumption: People educated in our school are more likely to do parallel programming at higher rather than lower tiers
- Language X10 not widely used in practice

Conclusions

- “Tiers of parallelism” is a fruitful concept
 - course structure
 - orientation for students
- Focus on lab sessions important
 - provided skeletons and solutions for every exercise
 - few students could chose their own language (typically much more complex than X10)
- X10 turned out to be very good choice
 - succinct expression of programs at different tiers
 - steep learning curve

Sources

- [1] Michael L. Scott: “Don’t start with Dekker’s algorithm – top-down introduction to concurrency”, Multicore Programming Education Workshop, 2009.
- [2] Michael L. Scott: “Making the simple case simple”, Position paper, Workshop on Curricula for Concurrency, in conjunction with OOPSLA, 2009.

Thank you for your attention.

Teaching materials are available at

<http://www.in.ohm-hochschule.de/professors/praun/pp>